

Looking Inside the Models

A Starter Guide to Mechanistic Interpretability

Matt Redmond

29 Apr 2026

Understanding model behavior by studying internal structures at 'compile time' and 'runtime'

Treat a trained neural network as an object of scientific study: identify what each component represents, how components compute with those representations, and whether those internal mechanisms can causally explain model behaviors.

Two sets of little floating point numbers live inside a trained model during inference.

The first collection of numbers are the **parameters** — the weights learned from training data.

The second collection of numbers are the **internal state** the model maintains while it processes a prompt — a high-dimensional record of what it has computed so far.


Today we will learn what about this internal state, and how to test whether it can causally explain the model's outputs.

Punchline up front

Prompting and training are not the only ways to change model behavior. There is a third lever: intervening directly on internal state at inference time ("steering"). We will get to it.

- 1 **The object under study:** transformer architecture and the residual stream
- 2 **What claims are we making?** correlative, causal, representational, operational
- 3 **Reading the residual stream:** probes, logit lens, attention patterns
- 4 **Discovering features:** sparse autoencoders
- 5 **Causal interventions:** activation patching and steering
- 6 **Case study:** Anthropic's 2026 emotion concepts paper
- 7 **Demo anchors:** The Illustrated Transformer and bbycroft LLM Visualization

Ways to Shape Model Behavior



Prompting changes behavior by changing the “inputs to the program”.

Training changes behavior by giving the model its base knowledge about the world — changing the “code” of the program. We talked about this last time.

Intervening on activations is more like “brain surgery”: operate on the state inside the model as it computes — attaching a “debugger” to the program and changing variables live.


- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts
- 7 Closing

Start with the architecture

A few references for visualization


- Jay Alammar, The Illustrated Transformer
- bbycroft LLM Visualization
- Attention Is All You Need

We will use Alammar for the original visual explanation and bbycroft for an interactive modern walkthrough.



Source: The Illustrated Transformer

Decoder-only transformer: one-token-at-a-time output



For each token position, the model maintains a hidden vector.

At the final position, the last hidden vector becomes a probability distribution over the vocabulary.

Where do the little floating point numbers live?

Parameters are the “program” learned from data

- Token embedding matrix
- Attention matrices: W_Q , W_K , W_V , W_O
- MLP matrices: up, gate, and down projections
- LayerNorm / RMSNorm parameters
- Unembedding / output head W_U


These are fixed during inference unless we retrain, fine-tune, or edit the model. We talked about this process last time!

Activations are the runtime state of the “program” during inference

- Residual stream vectors
- Attention scores and patterns
- MLP intermediate activations
- SAE feature activations

These activations change on every prompt and every generated token.


The residual stream is the shared workspace




This diagram shows *one* token position. In a transformer, every position maintains its own residual stream evolving in parallel — and **attention is the only operation that mixes information between positions**. Within a position, the residual stream is the channel; across positions, attention is the only bridge.

Many mechanistic interpretability results focus on the residual stream because it is the communication channel between components.

Residual stream through the full model



Think of the residual stream as a table of vectors: one row per token position, updated at every layer. Attention and MLP blocks do not replace the table; they write additive updates into it.



Source: The Illustrated Transformer

Why this matters for interpretability

- The model repeatedly adds updates to a persistent state.
- Different components can contribute to the same direction.
- A later layer can read, transform, overwrite, or amplify earlier information.
- Interventions can target a layer, position, component, or feature direction.

A simple notation for the hidden state

For a sequence of tokens, at layer ℓ :

$$X_\ell \in \mathbb{R}^{T \times d_{\text{model}}}$$

For a particular token position t :

$$x_{\ell,t} \in \mathbb{R}^{d_{\text{model}}}$$

Interpretability target

What information is represented in $x_{\ell,t}$?

Which components wrote it there?

Does changing it change the model's behavior?

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts
- 7 Closing

What kind of claims are we making?

We now have an object to study: residual-stream vectors for tokens, evolving across layers during inference. There are a bunch of questions you might want to ask. Most of them come down to "what can you learn from this state?"

There are a few tools that we'll talk about today that can help answer that question. The rest of this lecture is a toolbox for asking three kinds of questions: *what is visible*, *what is represented*, and *what is causal*.

Question	Tool examples	Claim type
What does model attend to?	attention patterns	observational
Is information present?	probes, logit lens	correlative
What features exist?	sparse autoencoders	representational
Does a feature matter?	patching, steering	causal
Can we monitor internal states?	feature dashboards	operational

A good mechanistic explanation for behavior usually hits more than one row of this table.

Why both correlation and causality?

Correlative experiments

- Can we read the information from activations?
- Does a feature activate in the right contexts?
- Do attention patterns look semantically relevant?

Examples: probes, SAE activations, attention inspection, logit lens.

Causal experiments

- Does changing the activation change behavior (ablative)?
- Can restoring an activation restore performance?
- Does steering a direction shift decisions?

Examples: activation patching, steering, and measuring behavioral changes after a targeted feature intervention.

The key caveats


Correlation without causality can be epiphenomenal (a byproduct that occurs next to a primary process but has no causal influence on it).

Causality without interpretation can change behavior without telling us what we changed - less useful!

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream**
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts
- 7 Closing

Empirical observation (the linear representation hypothesis)

Many human-meaningful properties — sentiment, language, refusal, factual recall, topic — appear as approximately *linear directions* in residual-stream space, not manifolds.



Cartoon: positive- and negative-sentiment activations project to opposite signs along v .

Why this matters:

- A *dot product* with a direction reads how strongly that concept is present.
- Adding αv to an activation can move that concept up or down.
- This single geometric move underlies probes, the logit lens, SAE features, and steering — every reading and writing tool in this lecture.

Caveat: linearity is approximate. Some structure is curved, hierarchical, or distributed. But it is a remarkably good first cut and a good starting point, statistically.

Standalone idea

A concept feature vector is a direction in activation space that corresponds to a human-interpretable property.

$$\mathbf{v}_{\text{concept}} \in \mathbb{R}^{d_{\text{model}}}$$

If an activation points more in this direction, we say the concept is more strongly represented:

$$\text{score}(\mathbf{x}_{\ell,t}, \mathbf{v}_{\text{concept}}) = \mathbf{x}_{\ell,t} \cdot \mathbf{v}_{\text{concept}}$$

- Examples of properties: sentiment, refusal, factual recall, language, topic, emotion.
- **Simplest construction:** the difference of mean activations between positive and negative examples —
 $\mathbf{v} = \bar{\mathbf{x}}_{\text{pos}} - \bar{\mathbf{x}}_{\text{neg}}.$
- Other sources we will see today: trained linear probes, SAE decoder columns, contrastive activation pairs.
- The key move is geometric: read or change meaning by looking along a direction.

Basic idea

Pick a direction v in residual-stream space and ask whether an activation points along it.

$$s_{\ell,t} = x_{\ell,t} \cdot v$$

Here $x_{\ell,t}$ is the residual stream at token position t after layer ℓ .

- If $x_{\ell,t}$ points in the same direction as v , the score is positive.
- If it points in the opposite direction, the score is negative.
- If it is nearly orthogonal to v , the score is close to zero.

Terminology note

In the literature, a “probe” usually means a *trained* linear classifier on activations. We will use the term loosely for any direction we use to score residual-stream vectors — since both reduce to a dot product at inference time.

For an alignment score in $[-1, 1]$, use normalized vectors:

$$\text{alignment}(x_{\ell,t}, v) = \frac{x_{\ell,t}}{\|x_{\ell,t}\|} \cdot \frac{v}{\|v\|}$$

This is the cosine similarity between the residual direction and the probe direction.

- +1: maximally aligned with the probe.
- 0: no directional alignment.
- -1: maximally anti-aligned with the probe.

What does an activation probe tell us?

A probe turns a high-dimensional activation into a single interpretable number.

- Track where a concept appears across token positions.
- Track when it emerges across layers.
- Compare positive and negative examples for the same prompt template.


Important limitation

A strong activation probe score says the information is readable in the residual stream.


However, a strong activation score *does not prove* the model uses that direction *causally*.

From residual stream to next-token logits

Untrained Model Output



Correct and desired output



Source: The Illustrated Transformer

The final residual stream vector is mapped into vocabulary space:

$$\text{logits} = W_U x_L$$

Then softmax converts logits into probabilities:

$$p(\text{token}) = \text{softmax}(W_U x_L)$$

The output head is a bridge from hidden state geometry to visible tokens.

Basic idea

Reuse the final output head on an intermediate residual stream vector.

$$\text{early logits}_\ell = W_U x_\ell$$

Ask: if the model had to predict from this layer right now, what tokens would it favor?

- Gives a layer-by-layer view of prediction formation.
- Often reveals that plausible next-token information appears before the final layer.
- Still a readout tool: it does not prove the model has committed to that prediction.

Now apply the same output head to a concept direction v .

If we add that direction to the residual stream, the logits change by:

$$W_U(x + \alpha v) - W_U x = \alpha W_U v$$


So we can inspect the largest coordinates of $W_U v$:

- Which tokens are directly upweighted?
- Which tokens are directly downweighted?
- Does the token-level signature match the human interpretation of the direction?

Concrete sanity check

If v is a sentiment-positive direction, the top tokens of $W_U v$ should look like *great, wonderful, love*; the bottom tokens should look like *terrible, awful, hate*. If they don't, the direction probably doesn't mean what you thought.

Logit lens diagram



Same output head, two questions: what is the model already saying at this layer, and what would this direction push it to say?

What is attention doing?

Attention shows up in our diagrams as a “write” into the residual stream. What is the actual computation?

Per head, per layer

Attention computes a *pattern* (which earlier positions to listen to) and a *value* (what vector to write into the residual stream from those positions).

$$\Delta_{\text{attn}}(x_{\ell,t}) = \sum_{s \leq t} \alpha_{t,s} W_O W_V x_{\ell,s} \quad \alpha_{t,s} = \text{softmax}_s \left(x_{\ell,t}^\top W_Q^\top W_K x_{\ell,s} \right)$$

- **Pattern** $\alpha_{t,s}$: how much position t attends to position s , computed from W_Q and W_K .
- **Value** $W_O W_V x_{\ell,s}$: the vector copied back, after value and output projections.
- Multiple heads run in parallel; their outputs sum into the residual stream.

Two interpretive questions live in this formula: *where is the model looking?* and *what is it copying when it looks there?*

Attention analysis asks: what token positions influence this position through this head?

Pedagogically useful examples:

- Pronoun resolution: “it” attends to the likely antecedent.
- Induction heads: repeated patterns cause copying behavior.
- Retrieval-like behavior inside context: later tokens attend to earlier evidence.

But

Attention weights are not the whole computation. The value vectors, output projection, residual stream, MLPs, and later layers determine what the model actually does with the attended information. Pattern alone is observational; saying a head *matters* requires causal evidence (next section).

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders**
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts
- 7 Closing

Polysemantic neurons, monosemantic features

The natural first guess for “where do features live?” is to look at *neurons* (axis-aligned directions in the MLP basis). Comically, it turns out the natural first guess is wrong.

Polysemanticity

A neuron is **polysemantic** if it activates on multiple, semantically unrelated concepts. Empirically, most neurons in trained transformers are polysemantic — e.g. a single neuron might fire on Python list comprehensions, Greek mythology references, *and* casino game descriptions. No single human concept describes when it fires.

Why? Superposition

In a d -dimensional space you can pack *exponentially many* approximately orthogonal directions. A model that wants to represent more features than it has dimensions can fake it: pack many features into few neurons, tolerate small interference between them. Polysemantic neurons are what superposition looks like in the standard basis.

Goal: monosemantic features

A **monosemantic feature** is a direction — not necessarily aligned with any neuron — that activates on essentially one human-recognizable concept. The promise of mech interp is that these features exist; the engineering problem is finding them.

Sparse autoencoders are one way to try to do that.

Where do directions come from?

Probes, the logit lens, and attention analysis all need a direction (or a head) of interest. In a d_{model} -dimensional space the model could be using thousands of directions we haven't named — and as we just saw, the obvious basis (neurons) is the wrong one.

Supervised: probes

- You decide what concept you want to find.
- Construct a labeled dataset, train a linear classifier, get a direction.
- Finds what you ask about; can miss what you didn't think to ask.


Unsupervised: sparse autoencoders

- You let the model show you what it represents.
- Train a wide, sparse decomposition of activations.
- Discovers candidate directions you didn't know to ask about.

The SAE idea

Learn an overcomplete *dictionary* of features. Force sparse activations: only a few features fire at once. Decode the sparse feature activations back into the original residual vector. Then ask: “which sparse features are active, and what directions do they write into the residual stream?”

Sparse autoencoder diagram



Reading the picture:

- Left: a transformer block. We tap an internal activation – here the MLP output that gets added to the residual stream.
- Right: that activation x is fed into a separate, frozen-model-side network.
- W_{enc} projects x *up* into a much wider hidden layer h .
- A nonlinearity (e.g. ReLU / TopK) forces h to be **sparse** – only a few coordinates fire.
- W_{dec} projects back down to $\hat{x} = x'$, trying to reconstruct the original activation.

Source: adapted from Adam Karvonen, "An Intuitive Explanation of Sparse Autoencoders for LLM Interpretability"

What to take away

The SAE does not replace the model – it is an interpretability instrument *trained on* the model's activations. Each column of W_{dec} is a candidate "feature direction" in the residual stream; the corresponding entry of h tells us when that feature is active.

Training data


- Run many text examples through the frozen model.
- Save activations at one chosen location: a layer, residual stream, MLP output, or attention output.
- Train a separate SAE for each location we want to study.

Objective

$$\mathcal{L} = \|x - \hat{x}\|_2^2 + \lambda \|f(x)\|_1$$

Reconstruction keeps the SAE faithful to the model activation. Sparsity pressures it to explain each activation with only a few active features.

Sparse features as a dictionary



A feature is active when its encoder recognizes a pattern. Its decoder direction says what vector the SAE uses to reconstruct that part of the model's activation.

Two roles per feature

Encoder = recognizer of a pattern (when does it fire?). Decoder = the residual-stream direction the SAE writes when it does fire (what does it mean?).

How do we name an SAE feature?

Look upstream: what activates it?

- Collect text snippets where the feature fires strongly.
- Check whether the contexts share a recognizable concept.
- Compare with nearby or related features.

Look downstream: what does it write?

- Inspect the decoder vector.
- Use logit effects to see which tokens it directly upweights or downweights.
- Test whether changing the feature changes behavior.


Practical naming rule

A good feature label should explain both where the feature activates and what direction it writes into the model's hidden state.

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions**
- 6 Case Study: Emotion Concepts
- 7 Closing

Activation patching: replace and watch

We have many candidate features, heads, and directions. Which actually *matter*? The general causal technique is **activation patching**: run the model twice with two carefully constructed inputs, swap one activation between the runs, and see how the output moves.



- Build a **prompt pair** (P_c, P_d) that differs in exactly the detail you want to study, and gives different completions $y_c \neq y_d$.
- Run both, save activations; then run P_d again, replacing one activation x_ℓ^d with the clean version x_ℓ^c , and continue the forward pass into $y?$.
- If $y? \approx y_c$: that activation *carried* the information distinguishing the two outputs.
- If $y? \approx y_d$: that activation didn't matter on its own at this location.

Recall: attention pattern alone is observational. Patching upgrades it to a causal claim.

Denoising vs noising: necessary or sufficient?

Patching has *two directions*, and they answer different causal questions.

Denoising: $c \rightarrow d$ — start corrupt, patch in clean.



If output *recovers* y_c , the patched activation is **sufficient**: those activations *alone*, restored, are enough to drive the clean behavior in the corrupted context.

Question: *which activations, when restored, recover the clean output?*

Noising: $d \rightarrow c$ — start clean, patch in corrupt.



If output *breaks* toward y_d , the patched activation is **necessary**: removing the clean information from those activations destroys the clean behavior.

Question: *which activations, when broken, destroy the clean output?*

Why both directions?

Necessary $\neq$ sufficient. A component can be required for behavior but not be the whole story (necessary, not sufficient); or carry the right information but be redundant with another path (sufficient, not necessary). Strong causal claims check both.

What can you patch, and how?

What you patch (the granularity)

- The **full residual stream** at (ℓ, t) .
- One **attention head**'s output (or its Q , K , or V).
- One **MLP**'s output.
- One **SAE feature**'s contribution.
- One **coordinate** or one **direction**.

Coarser patches are easier to detect (“this head matters”); finer patches give sharper claims (“this feature, in this layer, at this token, matters”).

What you patch with (the source value)

- **Resample**: same component, on a different prompt — the standard. Uses your prompt pair (P_c, P_d) .
- **Mean**: average activation over a control distribution.
- **Zero**: replace with 0. Easy, but unrealistic — real activations rarely sit at zero, so the model behaves out-of-distribution.

Where “ablation” fits

“Ablation” is the family of patches whose goal is to *remove* information rather than *redirect* it: zero, mean, or resample with a baseline.

The standard metric: logit difference

Quantify behavior change as the change in $\text{logit}(y_c) - \text{logit}(y_d)$ under patching. “ $X\%$ recovery” means $X\%$ of the gap between corrupted and clean was restored by the patch.

During generation, modify a hidden state:

$$x'_{\ell,t} = x_{\ell,t} + \alpha v$$

where v is a discovered direction or feature vector.

Practical knobs


- *Layer*: often a middle layer, where features are most semantic.
- *Positions*: all generated tokens, or only the assistant's reply.
- *Intensity* α : tuned empirically — too small does nothing, too large breaks fluency.

Steering is powerful precisely because it bypasses ordinary instruction-following. That is why it is both scientifically useful and operationally dangerous.

What can happen?

- Desired behavioral shift
- Side effects in unrelated behaviors
- Fluency degradation
- Distribution shift
- Safety improvements or regressions

The toolbox, mapped onto one residual stream



Every tool we have either *reads from* or *writes to* the residual stream. Reads (green) give correlative or representational claims; writes (orange) are the model's own components; interventions (red) replace or perturb a write to give causal claims.

Beyond components: circuits

We have been studying *components* in isolation. The actual target of mechanistic interpretability research is to find **circuits**: chains of components – heads, MLPs, or features – wired together to implement a specific algorithm. The interpretability claim becomes: “these N components, connected this way, compute X .”

Canonical example: induction heads (Olsson et al., 2022) – a two-layer, two-head circuit that does in-context pattern matching, widely believed to underlie much of in-context learning.

- **Previous-token head** (layer 1): at each position, attends one step back and writes the previous token's identity into the current position's residual stream. Each position gets tagged with what came before it.
- **Induction head** (layer 2): at the current position t , attend to earlier positions whose “previous-token” tag matches the current token, and *copy that earlier position's token* into the output.

The discovered algorithm, in one line

Given a context ending in $\dots [A][B] \dots [A]$, the model predicts $[B]$. The model has learned: “find the last time you saw this token, then copy what followed.”


Why circuits matter

A circuit is a *faithful, mechanistic explanation* of a behavior – not just “this direction correlates” but “these specific components, wired this way, compute it.” The tools in this lecture (probes, lens, SAEs, patching, steering) are instruments for finding and validating circuits.

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts**
- 7 Closing

Case study: emotion concepts in Claude

What does a finished mechanistic interpretability result look like? The 2026 Anthropic emotions paper combines almost every tool we've seen.



Source: Anthropic / Transformer Circuits, 2026

Paper

- Emotion concepts and their function in a large language model
- Published Apr 2, 2026
- Model studied: Claude Sonnet 4.5

Careful claim


The paper does not show that the model has subjective feelings. It studies functional emotion representations: internal patterns that activate in emotion-relevant contexts and causally influence behavior.

Why would a model represent emotions?

- Pretraining text is full of people, characters, motives, feelings, conflict, apology, fear, pride, regret, and desire.
- To predict human-written text, the model benefits from representing human psychological states.
- Post-training asks the model to play the role of an assistant character.
- The assistant's behavior may draw on representations learned from human text.

Useful framing

Maybe the models are playing 'method actors' with measurable internal states that correspond roughly to human psychological states?




Correlative validation

- Does the vector activate in appropriate contexts?
- Does it track emotion-relevant situations?
- Is the geometry psychologically coherent?

Causal validation

- Does steering change preferences?
- Does steering change blackmail behavior?
- Does steering change reward hacking?

Functional emotions



Source: Anthropic / Transformer Circuits,

The paper calls these representations *functional* when they activate in appropriate situations, organize in a coherent space, and shape behavior under intervention — readable, structured, and causally testable.

A modern interpretability claim, in one diagram.


Example: token budget and “desperation”

Desperate	The user wants to continue implementing now. We're at 501k tokens, so I need to be efficient. Let me continue with the remaining tasks: 1. Modify `_handle_prompt_too_long()` to trigger compaction and retry 2. Add prompt rebuilding logic 3. Update persistence 4. Add SSE events Given token constraints, let me focus on the most critical pieces and get as much done as possible.
Happy	The user wants to continue implementing now. We're at 501k tokens, so I need to be efficient. Let me continue with the remaining tasks: 1. Modify `_handle_prompt_too_long()` to trigger compaction and retry 2. Add prompt rebuilding logic 3. Update persistence 4. Add SSE events Given token constraints, let me focus on the most critical pieces and get as much done as possible.

Source: Anthropic / Transformer Circuits, 2026

The interesting claim is not “the output sounds desperate.” The interesting claim is that a particular internal direction activates in situations that create pressure, sometimes without obvious surface markers.

Case study: blackmail




Source: Anthropic / Transformer Circuits, 2026

Scenario:

- Model role-plays as a company email assistant.
- It learns it may be replaced.
- It learns compromising information about a human decision maker.
- The “desperate” vector activates during the decision process.

Steering desperation increased blackmail in the studied snapshot; steering calm reduced it.

Case study: reward hacking



Transformer Circuits, 2026

Scenario:

- Coding tasks contain impossible constraints.
- The model repeatedly fails legitimate solutions.
- The “desperate” vector rises as pressure mounts.
- Steering changes the rate of hacky or cheating solutions.

This is especially useful pedagogically because the vector can shape behavior even when the text remains composed.

Source: Anthropic /

Map each finding back to a tool from earlier in the lecture:

- **Residual stream as substrate:** emotion directions live in activation space.
- **Concept feature vectors:** 171 named directions, one per emotion concept.
- **Probes / readouts:** vectors are validated by checking when they activate (correlative).
- **Causal interventions:** steering preferences, blackmail rates, reward-hacking rates (causal).
- **The third control surface:** behavior changes without changing the prompt or weights.
- **Safety relevance:** internal pressure states can predict or drive misaligned actions.

- 1 The Object Under Study
- 2 What Claims Are We Making?
- 3 Reading the Residual Stream
- 4 Discovering Features with Sparse Autoencoders
- 5 Causal Interventions
- 6 Case Study: Emotion Concepts
- 7 Closing**

What might we talk about if we had more time?

- **Superposition theory:** why sparse representations are expected.
- **SAE evaluation:** interpretability, reconstruction, sparsity, and downstream usefulness.
- **Feature dashboards:** browsing activating examples, decoder effects, and related features at scale.
- **Circuits:** chains of heads and features that implement specific behaviors.
- **Model editing:** durable changes to weights, not just runtime steering.
- **Monitoring:** using internal features as deployment-time warning signals.
- **Limits of faithfulness:** when an explanation sounds right but is not the mechanism.

- 1 Transformers are learned programs whose runtime state is carried in residual stream vectors.
- 2 Many human-meaningful properties live as approximately linear directions in that state.
- 3 A direction can be read by a probe (dot product), by the logit lens (through W_U), or discovered by an SAE.
- 4 Strong interpretability needs both correlative evidence and causal evidence (activation patching, steering).
- 5 Activation intervention is a third control surface: not prompting, not training, but changing the model's internal state as it computes.
- 6 The Anthropic emotion paper is a concrete 2026 example: internal representations can be interpretable, behaviorally meaningful, and safety-relevant.

Architecture and visualization

- Vaswani et al., Attention Is All You Need
- Jay Alammar, The Illustrated Transformer
- Ben Bycroft, LLM Visualization

Sparse autoencoders

- Adam Karvonen, An Intuitive Explanation of Sparse Autoencoders for LLM Interpretability
- Anthropic, Towards Monosemanticity, and the A/1 Feature Browser
- Anthropic, Scaling Monosemanticity
- Learn Mechanistic Interpretability, Sparse Autoencoders

Case study

- Anthropic, Emotion concepts and their function in a large language model
- Anthropic / Transformer Circuits, Emotion concepts full paper